

Computer Systems II

Creating and Executing Processes

Unix system calls

fork()


wait()

exit()

How To Create New Processes?


■ Underlying mechanism

- A process runs **fork** to create a child process
- Parent and children execute concurrently
- Child process is a duplicate of the parent process



Process Creation

- After a **fork**, both parent and child keep running, and each can fork off other processes.
- A **process tree** results. The root of the tree is a special process created by the OS during startup.
- A process can *choose* to wait for children to terminate. For example, if C issued a **wait()** system call, it would block until G finished.

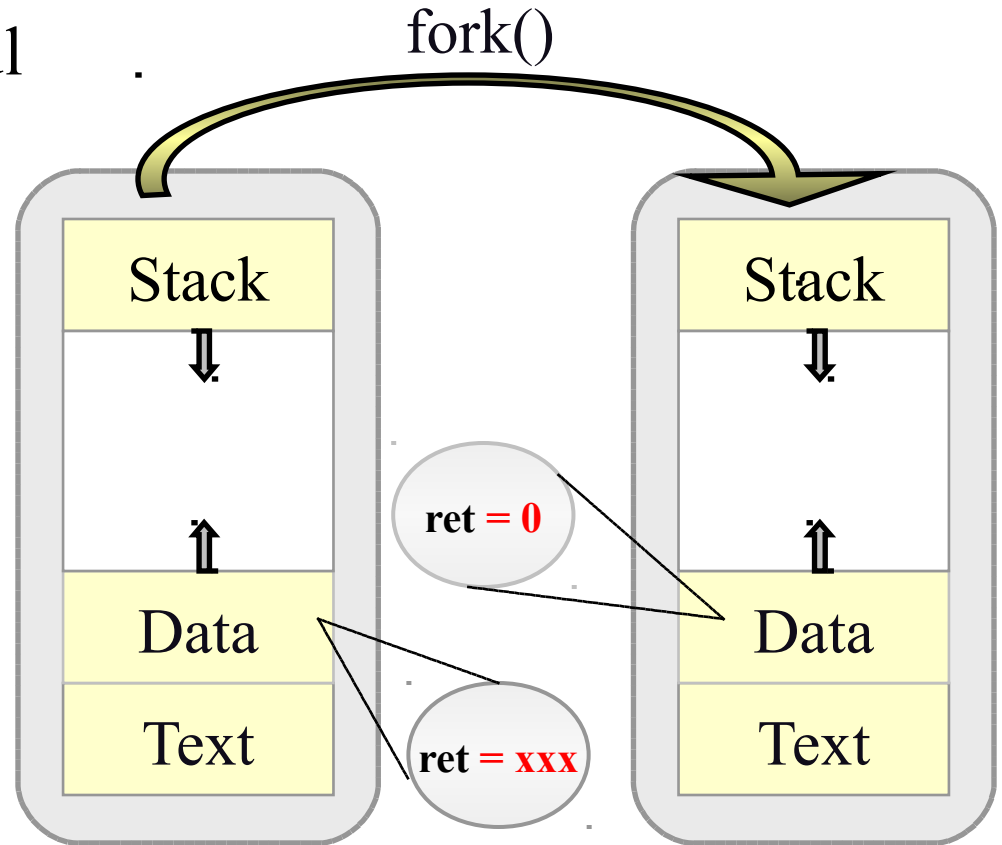


Bootstrapping

- When a computer is switched on or reset, there must be an initial program that gets the system running
- This is the bootstrap program
 - Initialize CPU registers, device controllers, memory
 - Load the OS into memory
 - Start the OS running
- OS starts the first process (such as “init”)
- OS waits for some event to occur
 - Hardware interrupts or software interrupts (traps)

Fork System Call

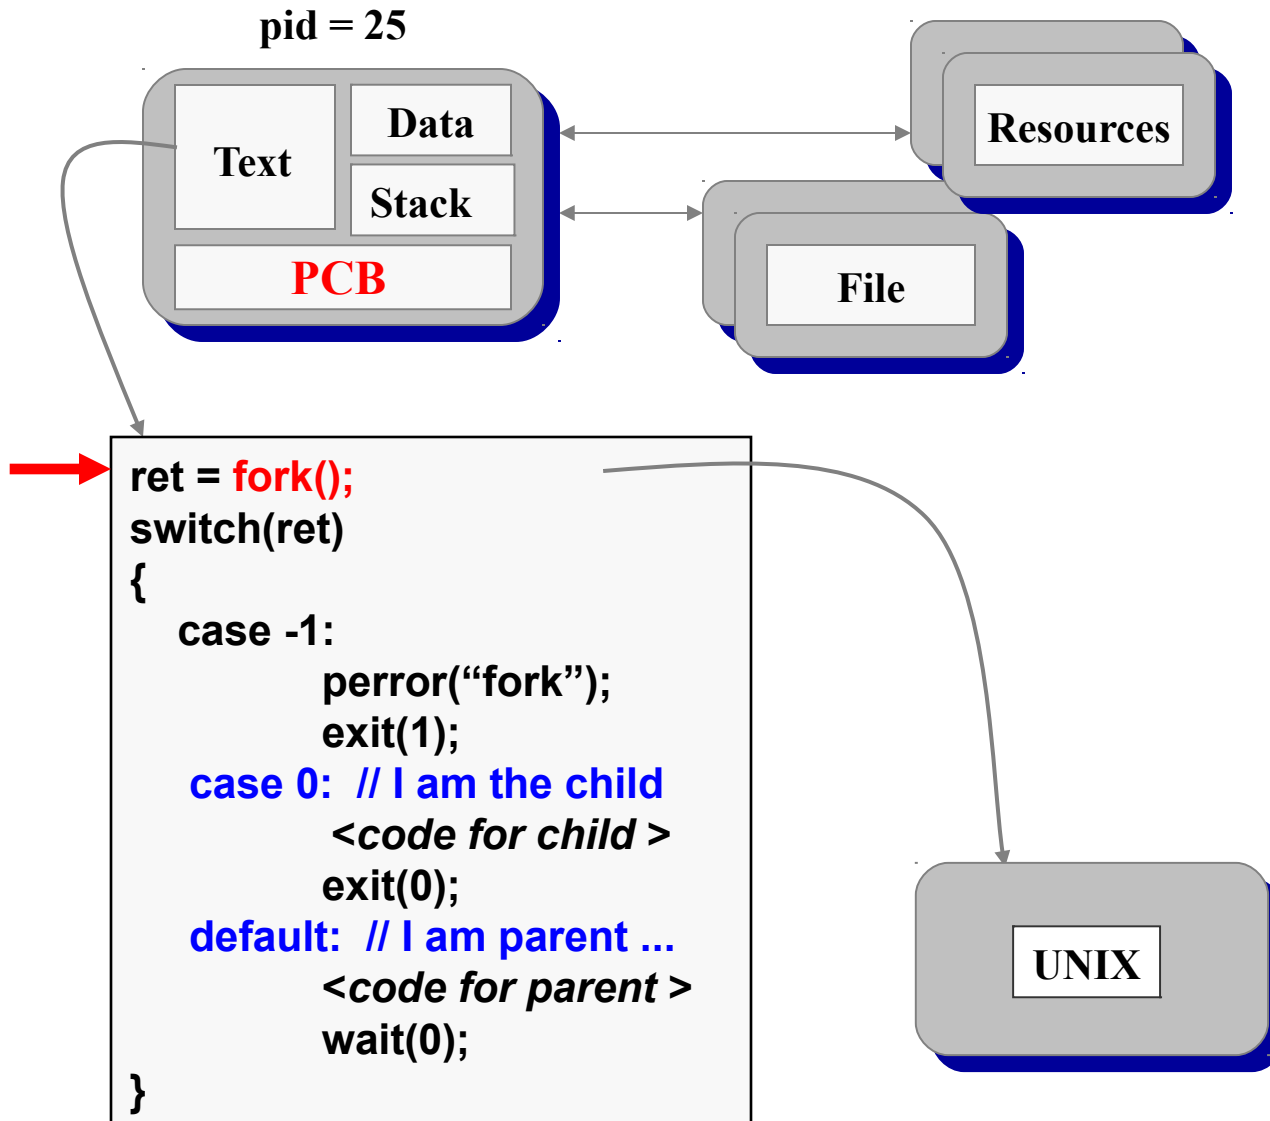
- Current process split into 2 processes: parent, child
- Returns -1 if unsuccessful
- Returns 0 in the child
- Returns the child's identifier in the parent



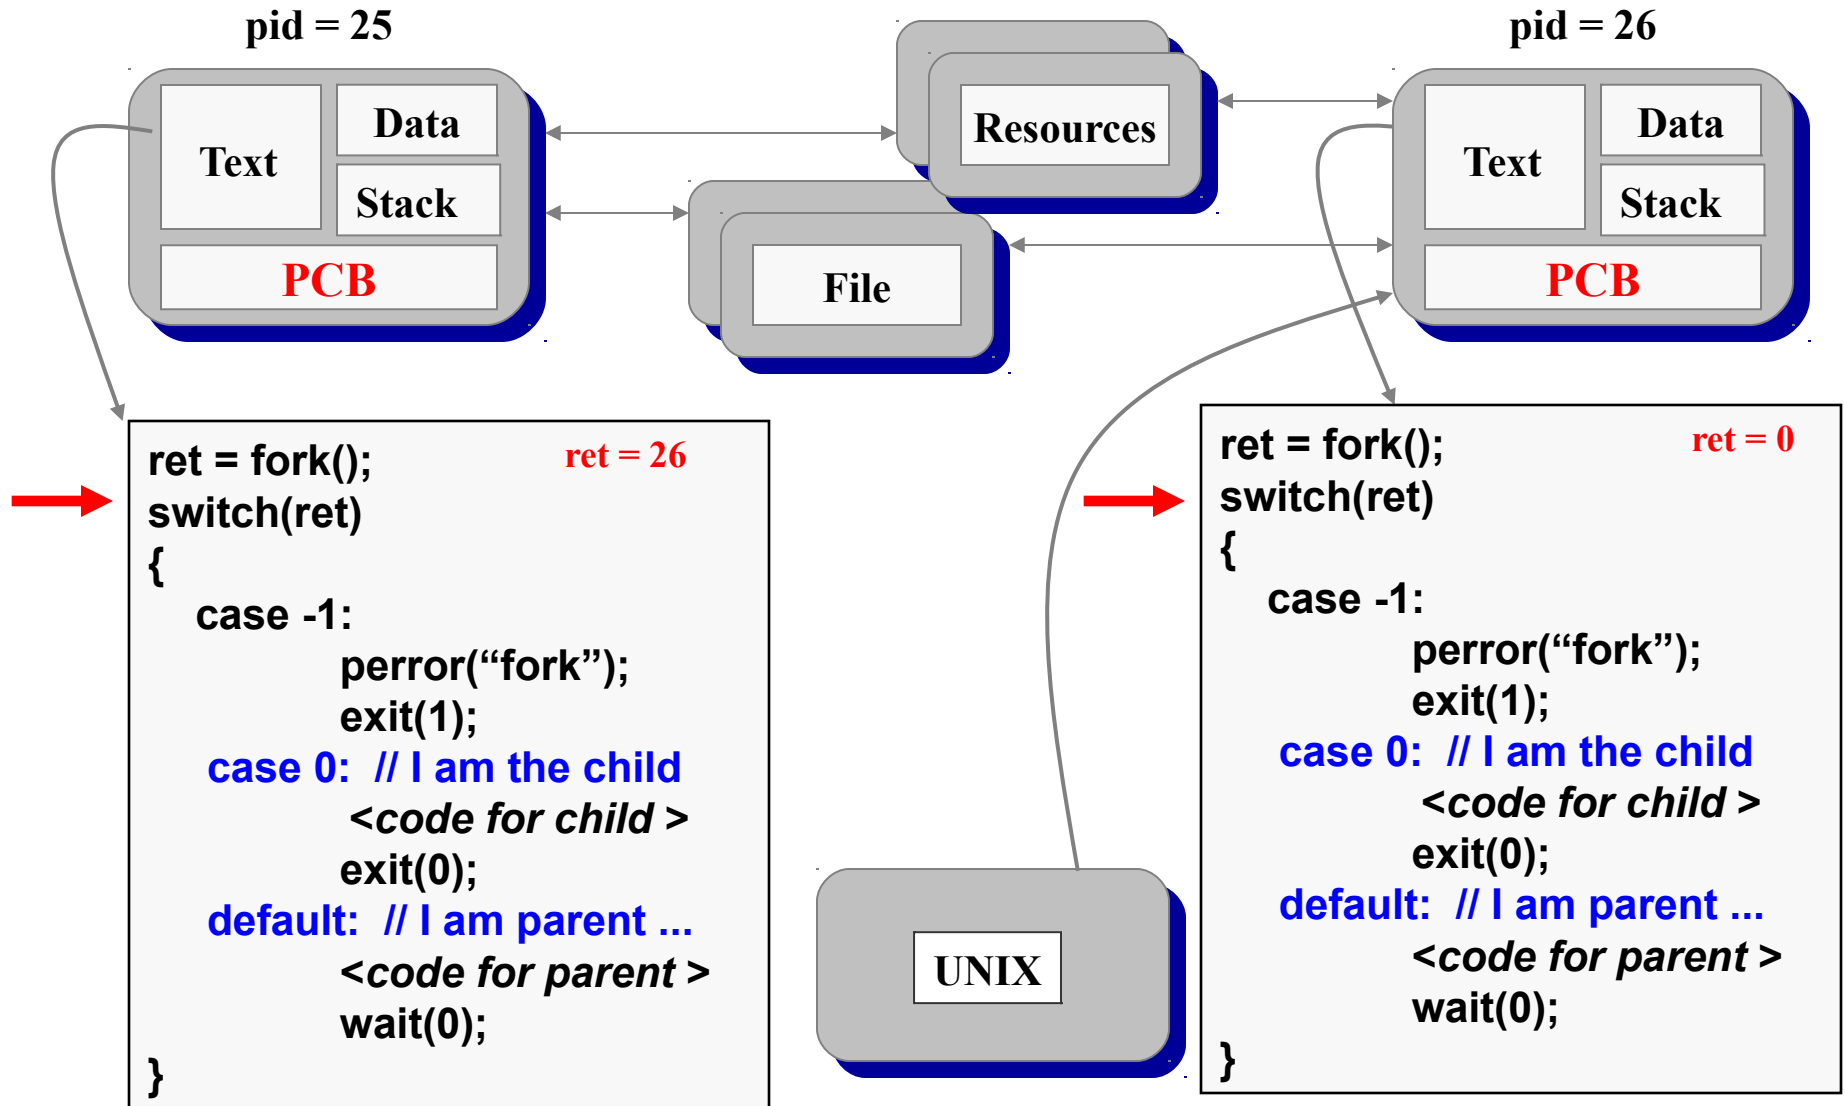
Fork System Call

- The child process inherits from parent
 - identical copy of memory
 - CPU registers
 - all files that have been opened by the parent
- Execution proceeds **concurrently** with the instruction following the fork system call
- The execution context (PCB) for the child process is a copy of the parent's context at the time of the call

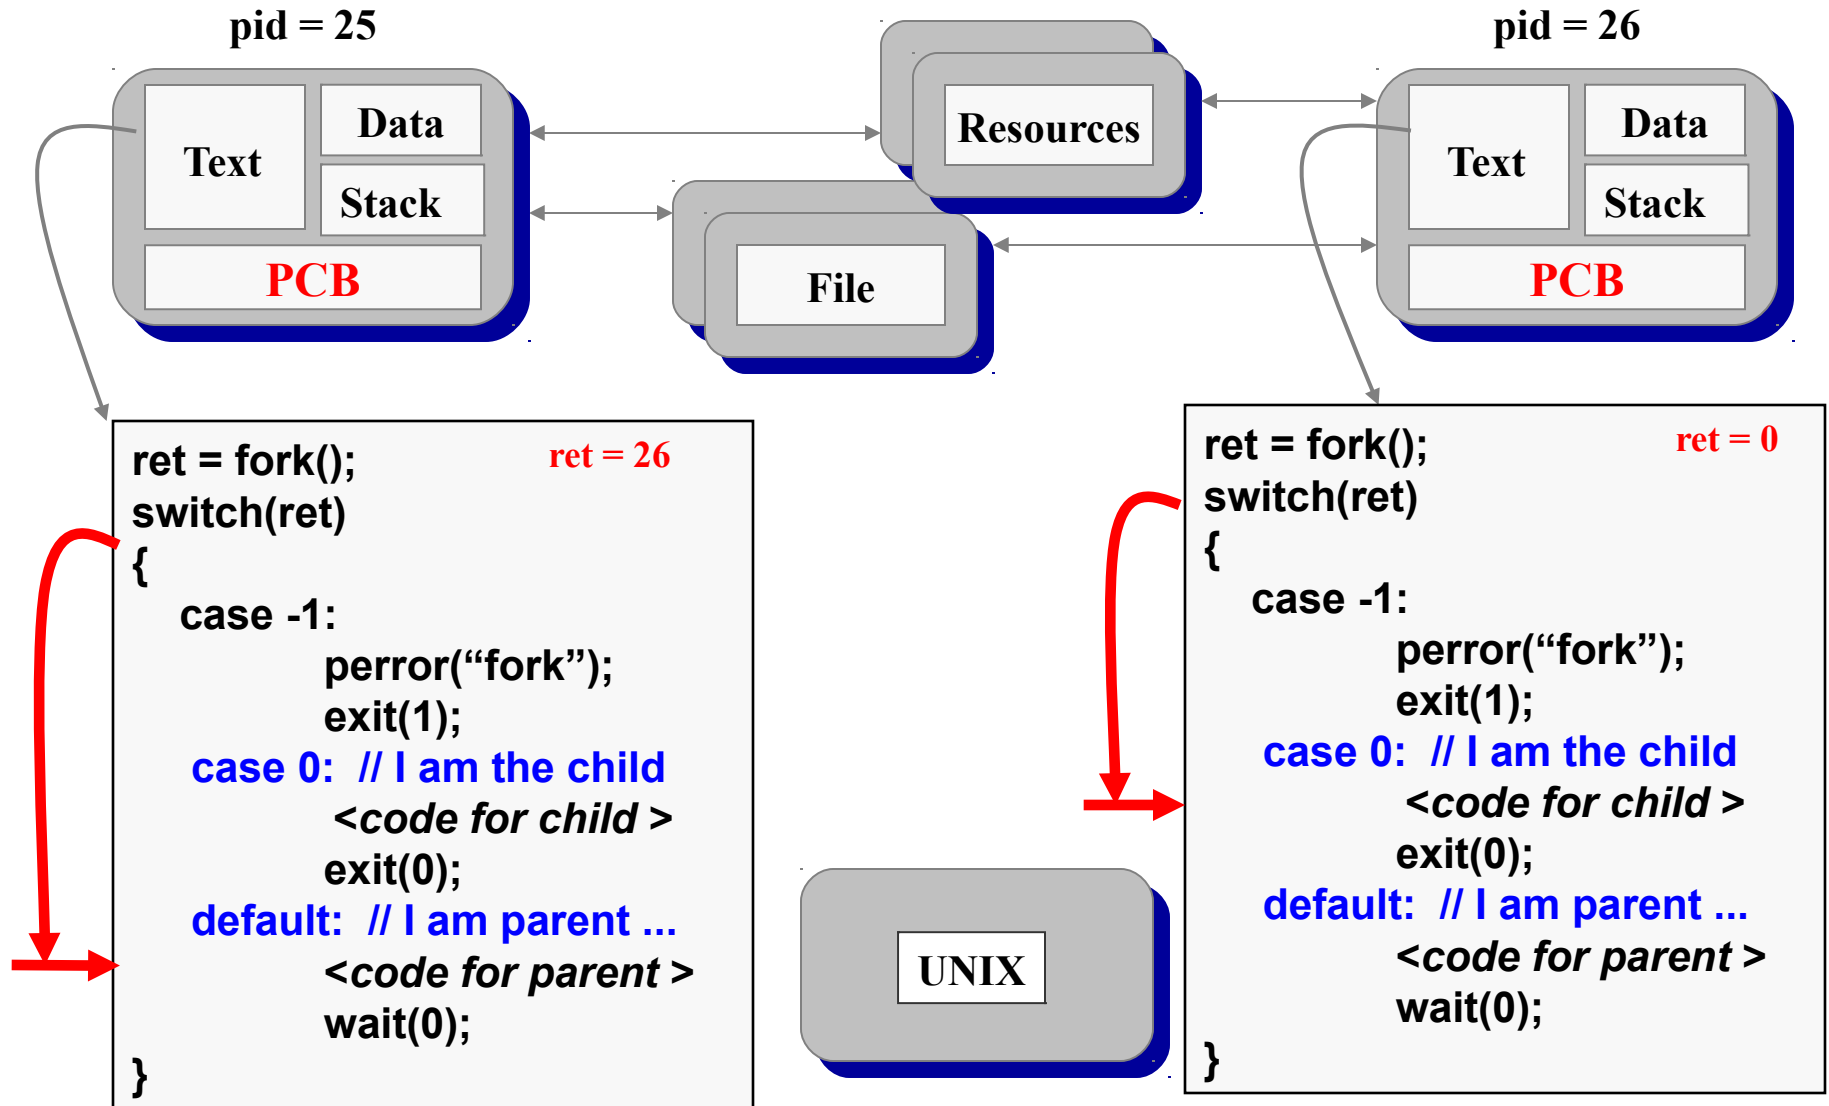
How fork Works (1)



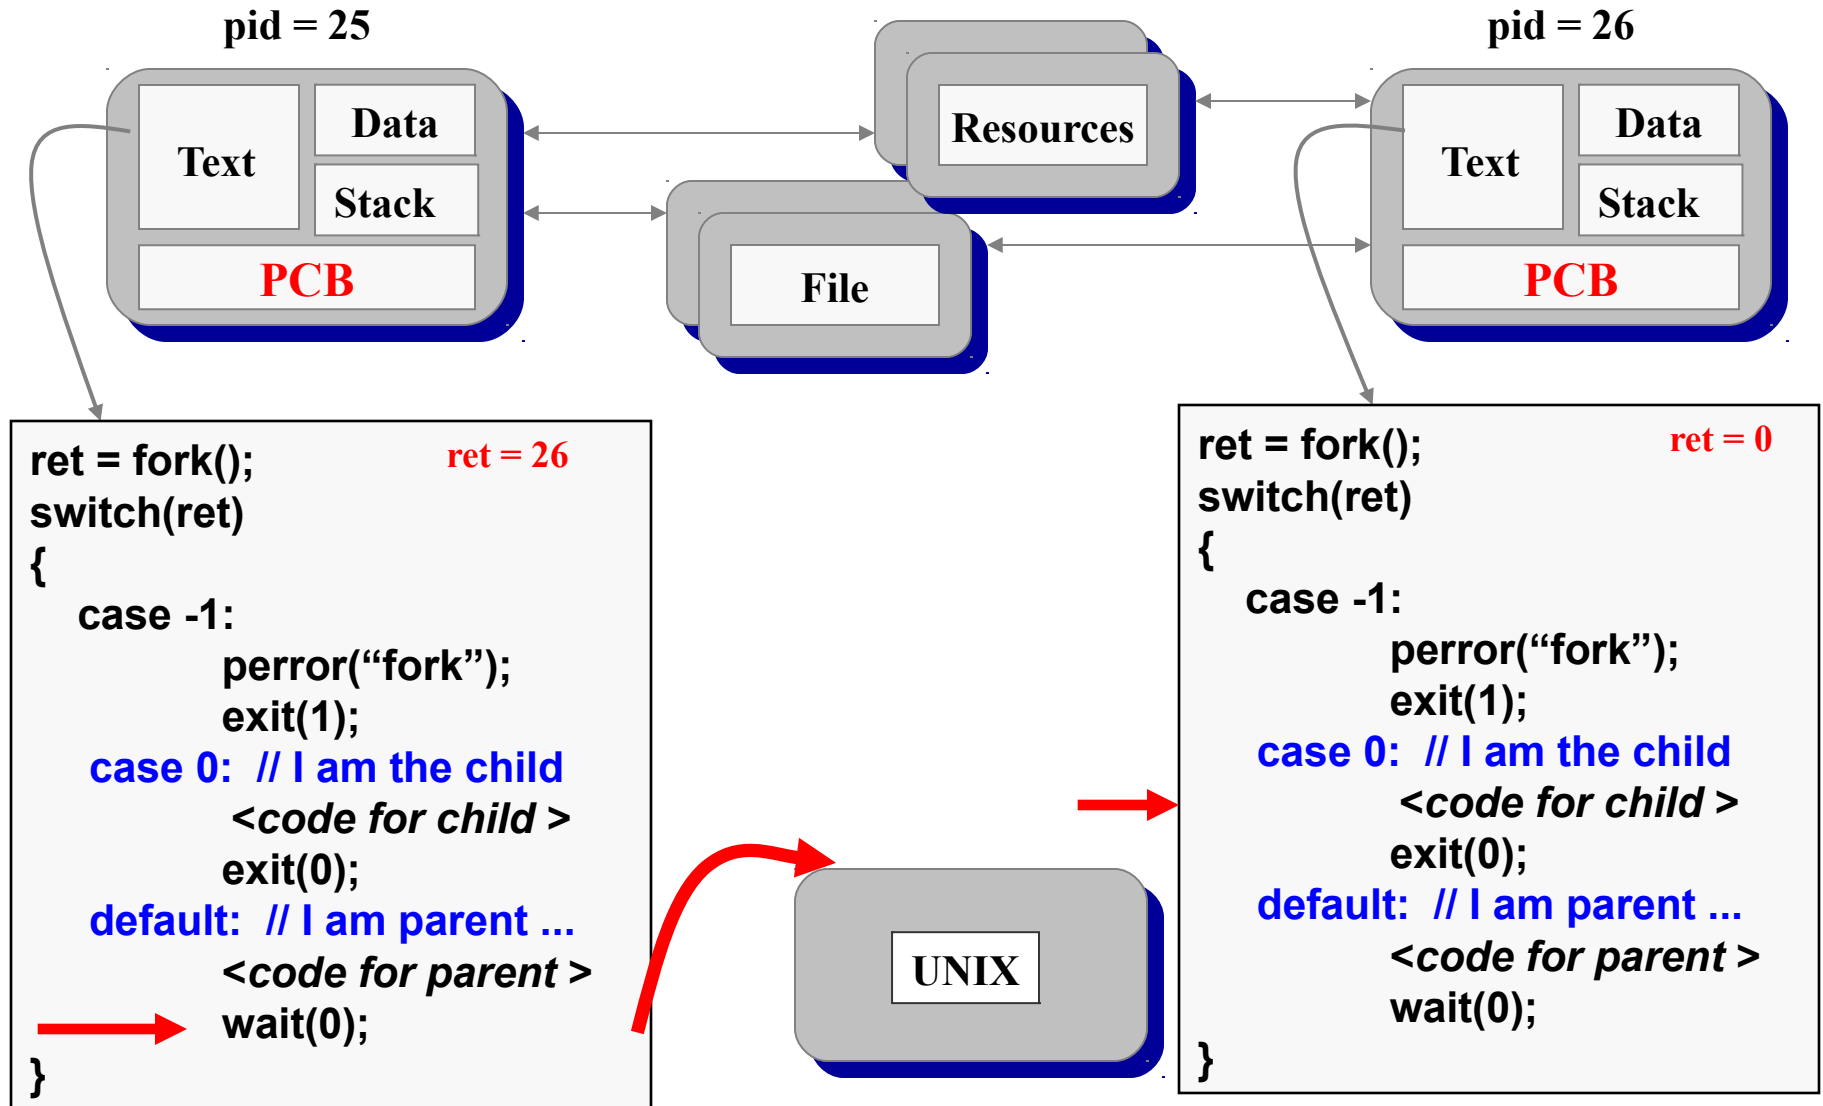
How fork Works (2)



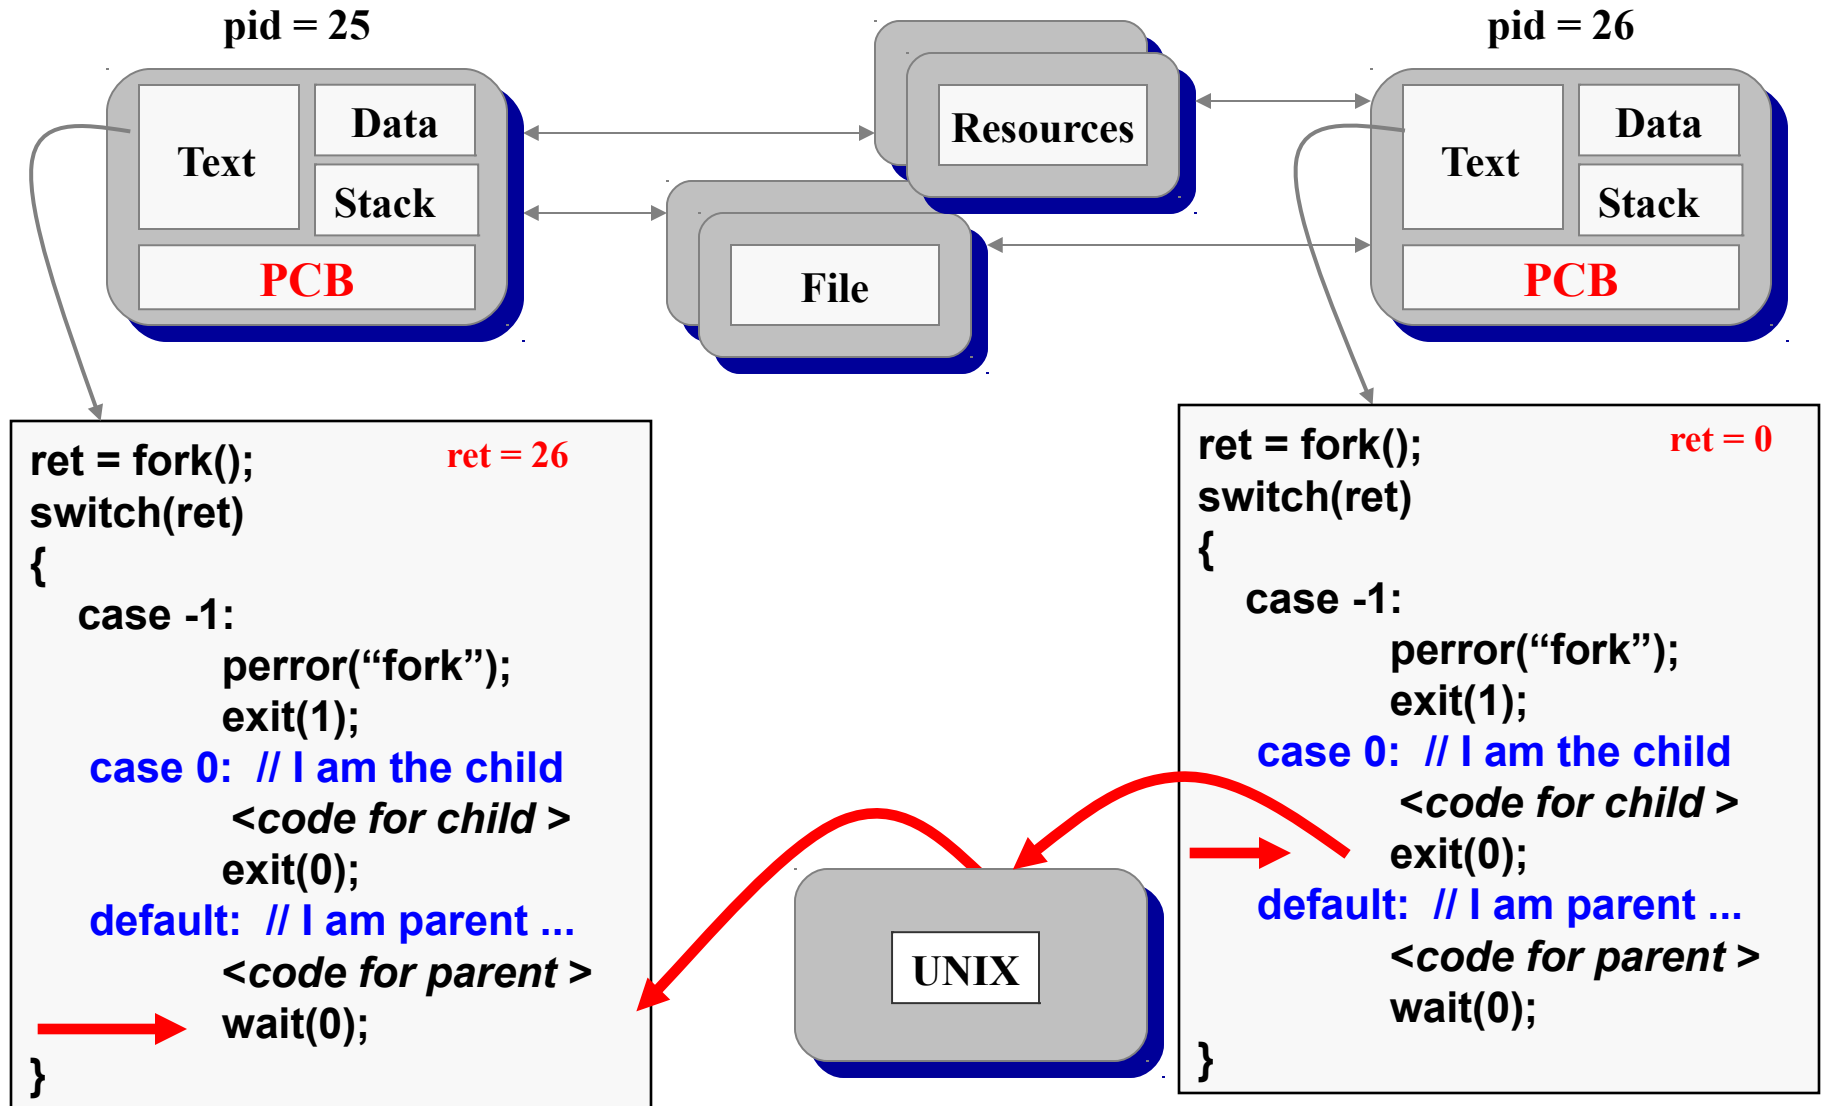
How fork Works (3)



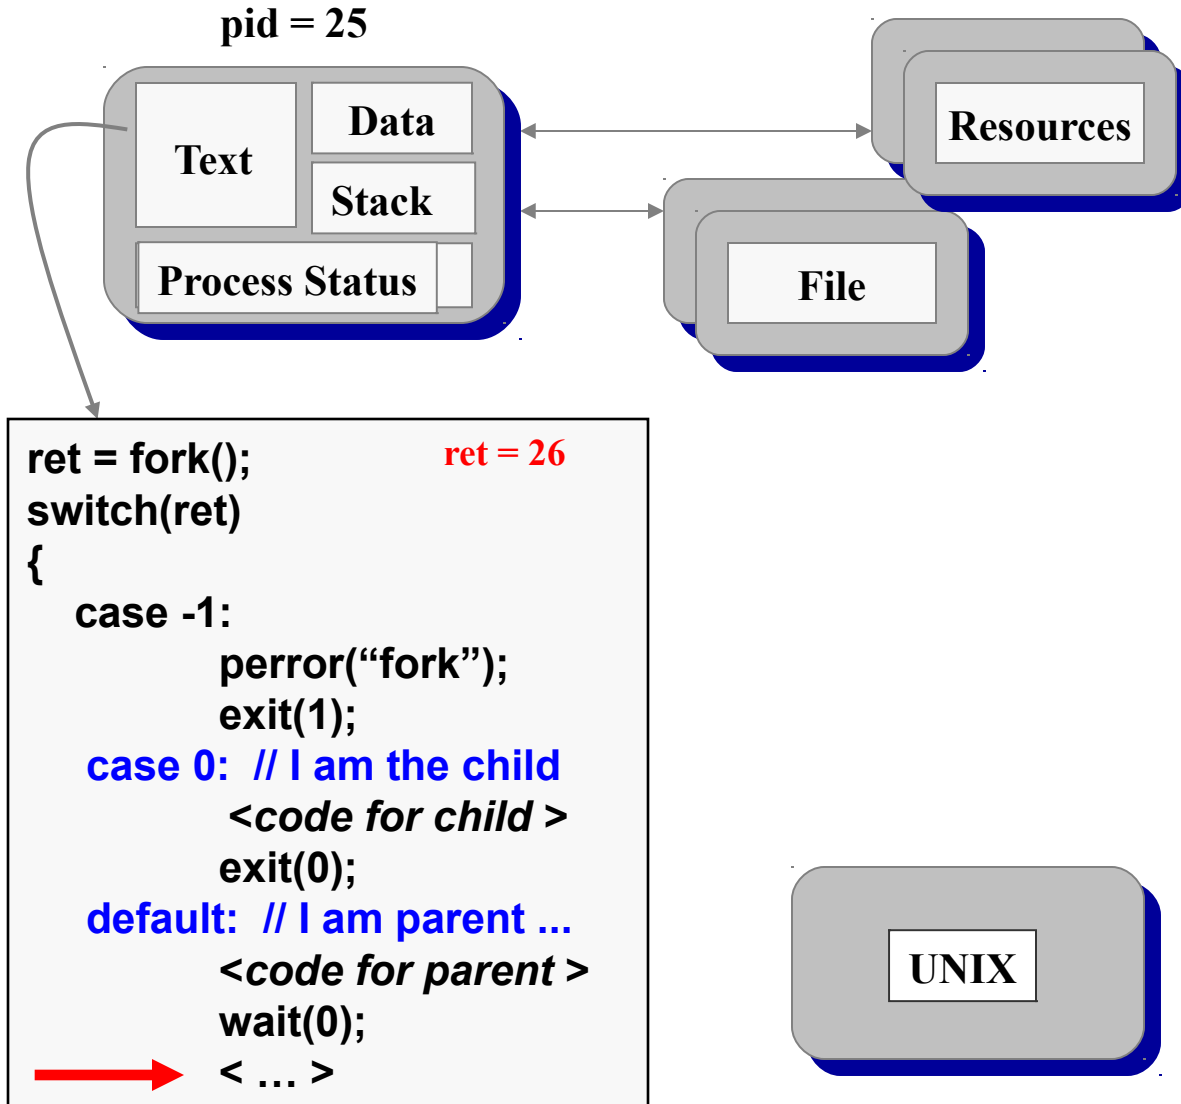
How fork Works (4)



How fork Works (5)



How fork Works (6)



Orderly Termination: `exit()`

- To finish execution, a child may call `exit(number)`
- This system call:
 - Saves result = argument of `exit`
 - Closes all open files, connections
 - Deallocates memory
 - Checks if parent is alive
 - If parent is alive, holds the result value until the parent requests it (with `wait`); in this case, the child process does not really die, but it enters a zombie/defunct state
 - If parent is not alive, the child terminates (dies)


Waiting for the Child to Finish

- Parent may want to wait for children to finish
 - Example: a shell waiting for operations to complete
- Waiting for any some child to terminate: `wait()`
 - Blocks until some child terminates
 - Returns the process ID of the child process
 - Or returns -1 if no children exist (i.e., already exited)
- Waiting for a specific child to terminate: `waitpid()`
 - Blocks till a child with particular process ID terminates

```
#include <sys/types.h>
#include <sys/wait.h>

pid_t wait(int *status);
pid_t waitpid(pid_t pid, int *status, int options);
```

State Transition on wait and exit Calls



Other useful system calls: `getpid`, `getppid`

- `getpid` returns the identifier of the calling process. Example call (`pid` is an integer):

```
pid = getpid();
```

- `getppid` returns the identifier of the parent.
- Note:
 - Zombies can be noticed by running the '`ps`' command (shows the process list); you will see the string "<`defunct`>" as their command name:

```
ps -ef
```

```
ps -ef | grep mdamian
```

Fork Example 1: What Output?

```
int main()
{
    pid_t pid;
    int x = 1;


    pid = fork();
    if (pid != 0) {
        printf("parent: x = %d\n", --x);
        exit(0);
    } else {
        printf("child: x = %d\n", ++x);
        exit(0);
    }
}
```

Fork Example 2

■ Key Points


- Both parent and child can continue forking

```
void fork2()  
{  
    printf("L0\n");  
    fork();  
    printf("L1\n");  
    fork();  
    printf("Bye\n");  
}
```



Fork Example 3

```
void fork3()  
{  
    printf("L0\n");  
    fork();  
    printf("L1\n");  
    fork();  
    printf("L2\n");  
    fork();  
    printf("Bye\n");  
}
```



Fork Example 4

```
void fork4()  
{  
    printf("L0\n");  
    if (fork() != 0) {  
        printf("L1\n");  
        if (fork() != 0) {  
            printf("L2\n");  
            fork();  
        }  
    }  
    printf("Bye\n");  
}
```

Fork Example 5

```
void fork5()  
{  
    printf("L0\n");  
    if (fork() == 0) {  
        printf("L1\n");  
        if (fork() == 0) {  
            printf("L2\n");  
            fork();  
        }  
    }  
    printf("Bye\n");  
}
```

Summary

■ Fork

- Creates a duplicate of the calling process
- The result is two processes: parent and child
- Both continue executing from the same point on

■ Exit

- Orderly program termination
- Unblocks waiting parent

■ Wait

- Used by parent
- Waits for child to finish execution

Unix system calls

execv

execl

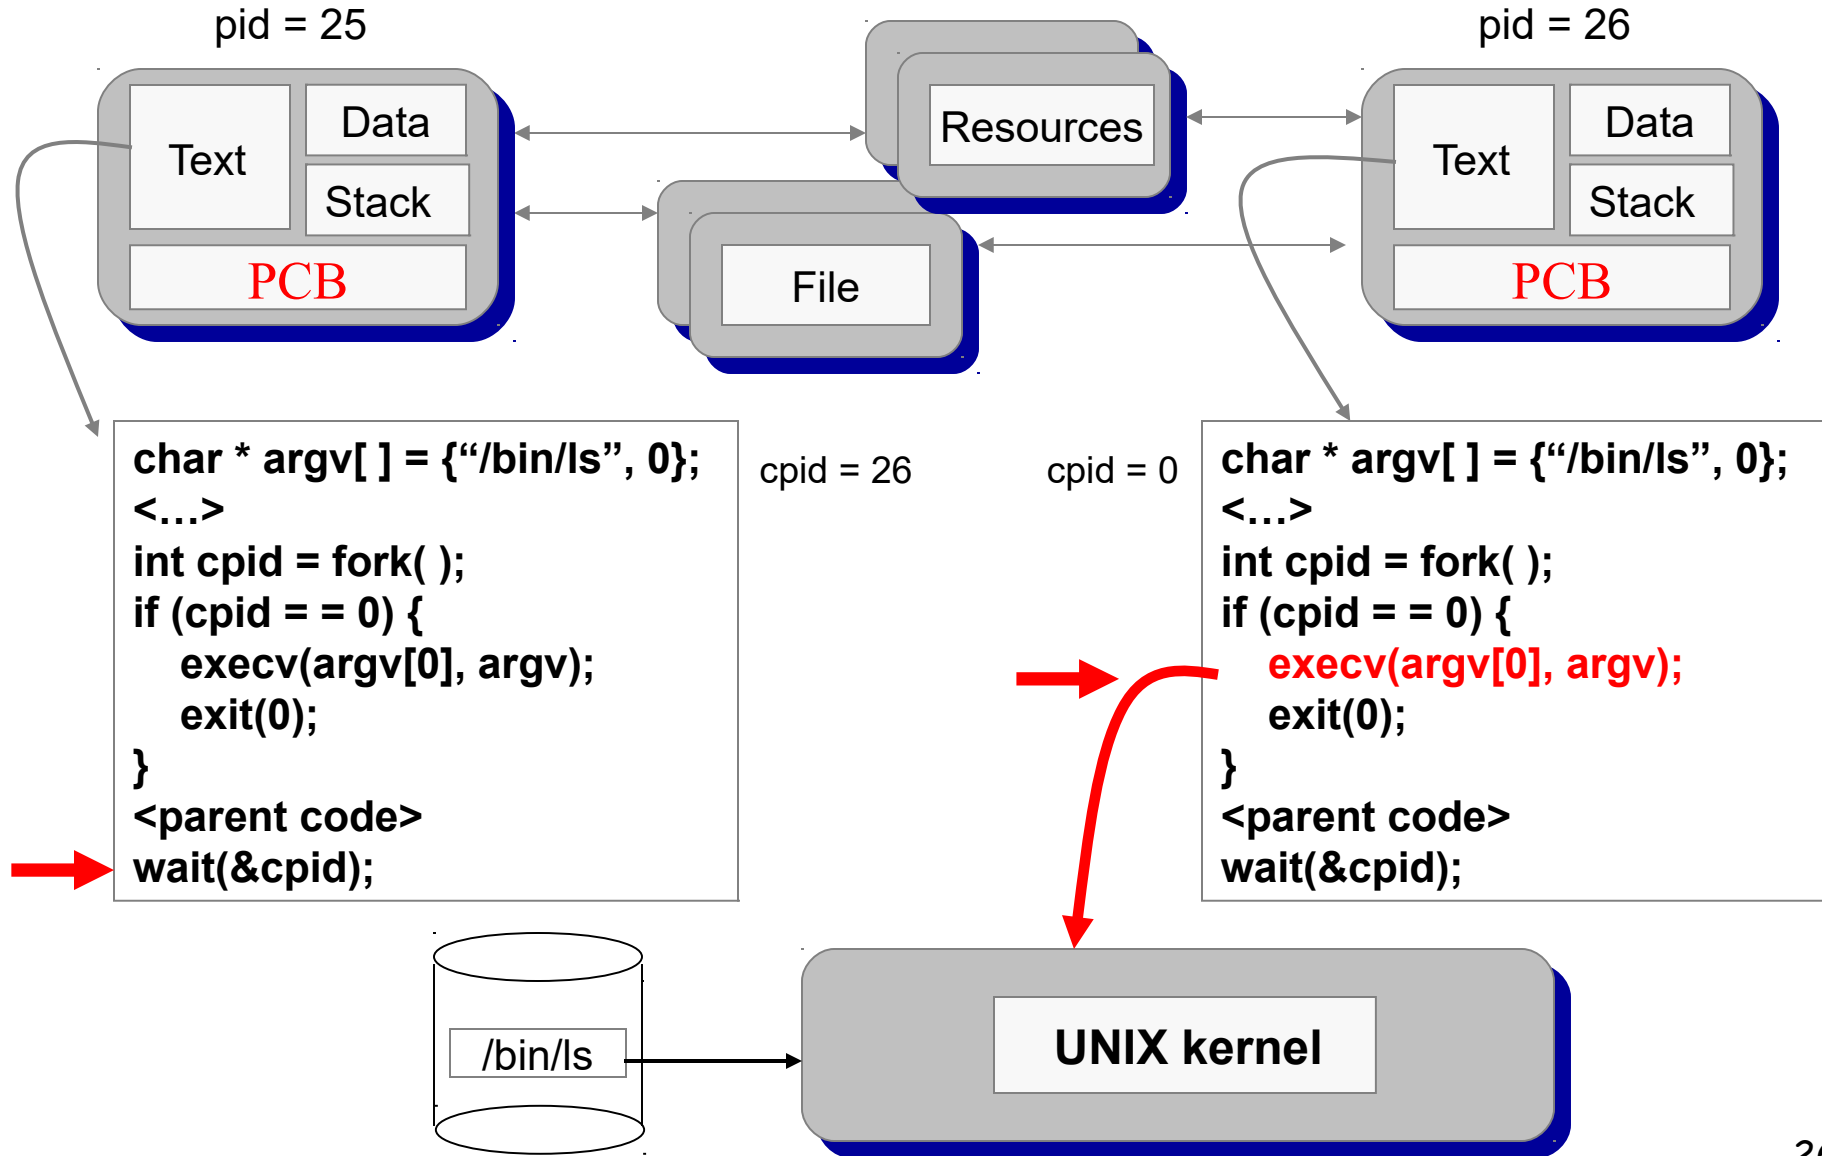
Unix's `execv`

- The system call `execv` executes a file, transforming the calling process into a new process. After a successful `execv`, there is no return to the calling process.

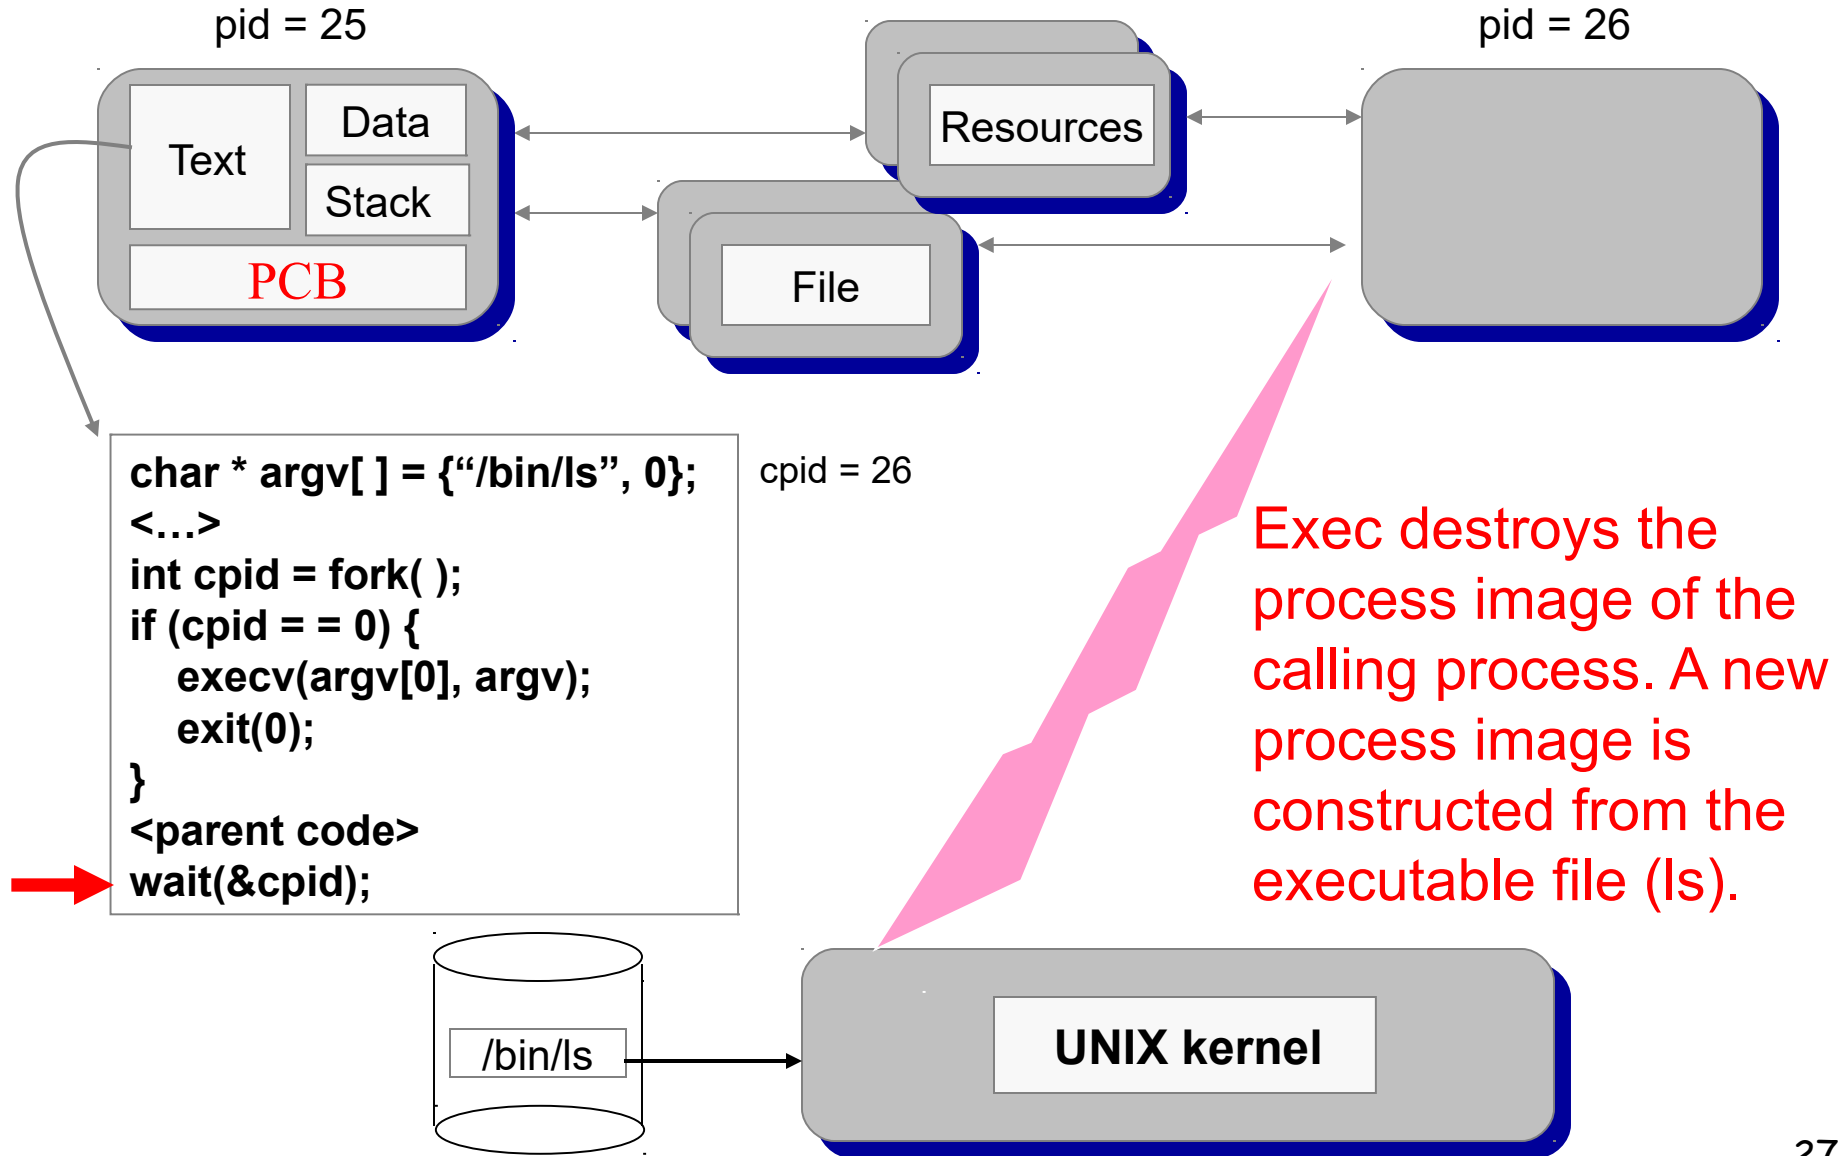
```
execv(const char * path, char * const argv[])
```

- `path` is the full path for the file to be executed
- `argv` is the array of arguments for the program to execute
 - each argument is a null-terminated string
 - the first argument is the name of the program
 - the last entry in `argv` is `NULL`

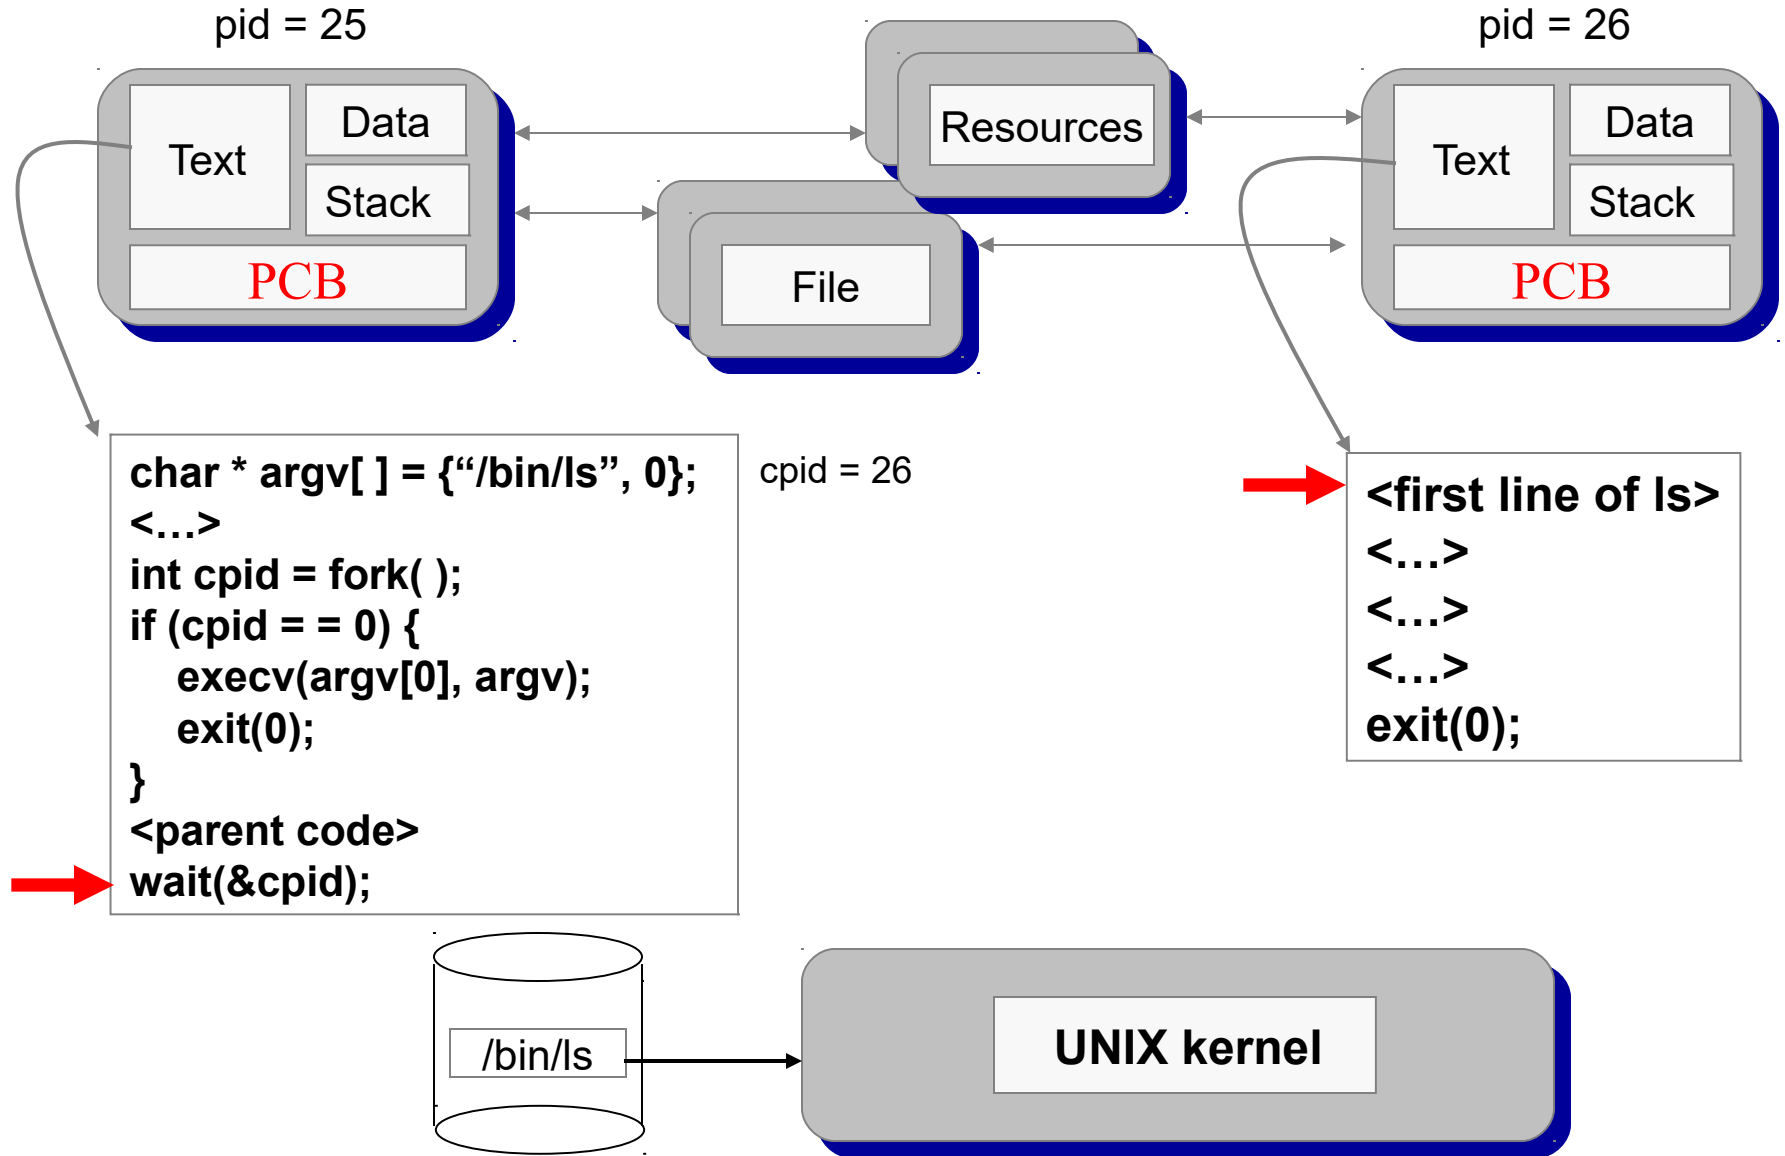
How execv Works (1)



How execv Works (2)




How execv Works (3)



Example: A Simple Shell

- Shell is the parent process
 - E.g., bash
- Parses command line
 - E.g., “ls -l”
- Invokes child process
 - Fork, execvp
- Waits for child
 - Wait




execv Example

```
#include <stdio.h>
#include <unistd.h>

char * argv[] = {"/bin/ls", "-l", 0};
int main()
{
    int pid, status;

    if ( (pid = fork() ) < 0 )
    {
        printf("Fork error \n");
        exit(1);
    }
    if(pid == 0) { /* Child executes here */
        execv(argv[0], argv);
        printf("Exec error \n");
        exit(1);
    } else /* Parent executes here */
        wait(&status);
    printf("Hello there! \n");
    return 0;
}
```



Note the NULL string
at the end

execv Example – Sample Output

- Sample output:

total 282

drwxr-xr-x 2 mdamian faculty 512 Jan 29 14:02 assignments

-rw-r--r-- 1 mdamian faculty 3404 Jan 29 14:05 index.html

drwxr-xr-x 2 mdamian faculty 512 Jan 28 15:02 notes

Hello there!

execl

- Same as execv, but takes the arguments of the new program as a list, not a vector:

- Example:

```
execl("/bin/ls", "/bin/ls", "-l", 0);
```

- Is equivalent to

Note the NULL string at the end

```
char * argv[] = {"/bin/ls", "-l", 0};  
execv(argv[0], argv);
```

- execl is mainly used when the number of arguments is known in advance

General purpose process creation

- In the parent process:

```
int childPid;
char * const argv[ ] = {...};

main {
    childPid = fork();
    if(childPid == 0)
    {
        // I am child ...
        // Do some cleaning, close files
        execv(argv[0], argv);
    }
    else
    {
        // I am parent ...
        <code for parent process>
        wait(0);
    }
}
```

Combined fork/exec/wait

- Common combination of operations
 - Fork to create a new child process
 - Exec to invoke new program in child process
 - Wait in the parent process for the child to complete
- Single call that combines all three
 - `int system(const char *cmd);`
- Example

```
int main()  
{  
    system("echo Hello world");  
}
```

Properties of fork / exec sequence

- In 99% of the time, we call `execv(...)` after `fork()`
 - the memory copying during `fork()` is useless
 - the child process will likely close open files and connections
 - overhead is therefore high
 - might as well combine both in one call (`OS/2`)
- `vfork()`
 - a system call that creates a process without creating an identical memory image
 - sometimes called “lightweight” fork
 - child process is understood to call `execv()` almost immediately

Variations of `execv`

- `execv`

- Program arguments passed as an array of strings

- `execvp`

- Extension of `execv`
- Searches for the program name in the `PATH` environment

- `execl`

- Program arguments passed directly as a list

- `execlp`

- Extension of `execv`
- Searches for the program name in the `PATH` environment

Summary

■ `exec(v,vp,l,lp)`

- Does NOT create a new process
- Loads a new program in the image of the calling process
- The first argument is the program name (or full path)
- Program arguments are passed as a vector (`v,vp`) or list (`l,lp`)
- Commonly called by a forked child

■ `system:`

- combines fork, wait, and exec all in one