

# Signals

# Context of Remaining Lectures



Second half of COS 217 takes 2 tours:

1. “Language levels” tour

**C → assembly language → machine language**

Illustrated by assembly language asgt, buffer overrun asgt

2. “Service levels” tour

**C → C standard library → operating system (OS)**

Illustrated by heap manager asgt, shell asgt

The 2 remaining lectures flesh out the “service levels” tour



# Goals of Remaining Lectures

Two fundamental questions:

Q1: How does the **OS** communicate to an **application process**?

Q2: How does an **application process** communicate to the **OS**?

This lecture: Q1

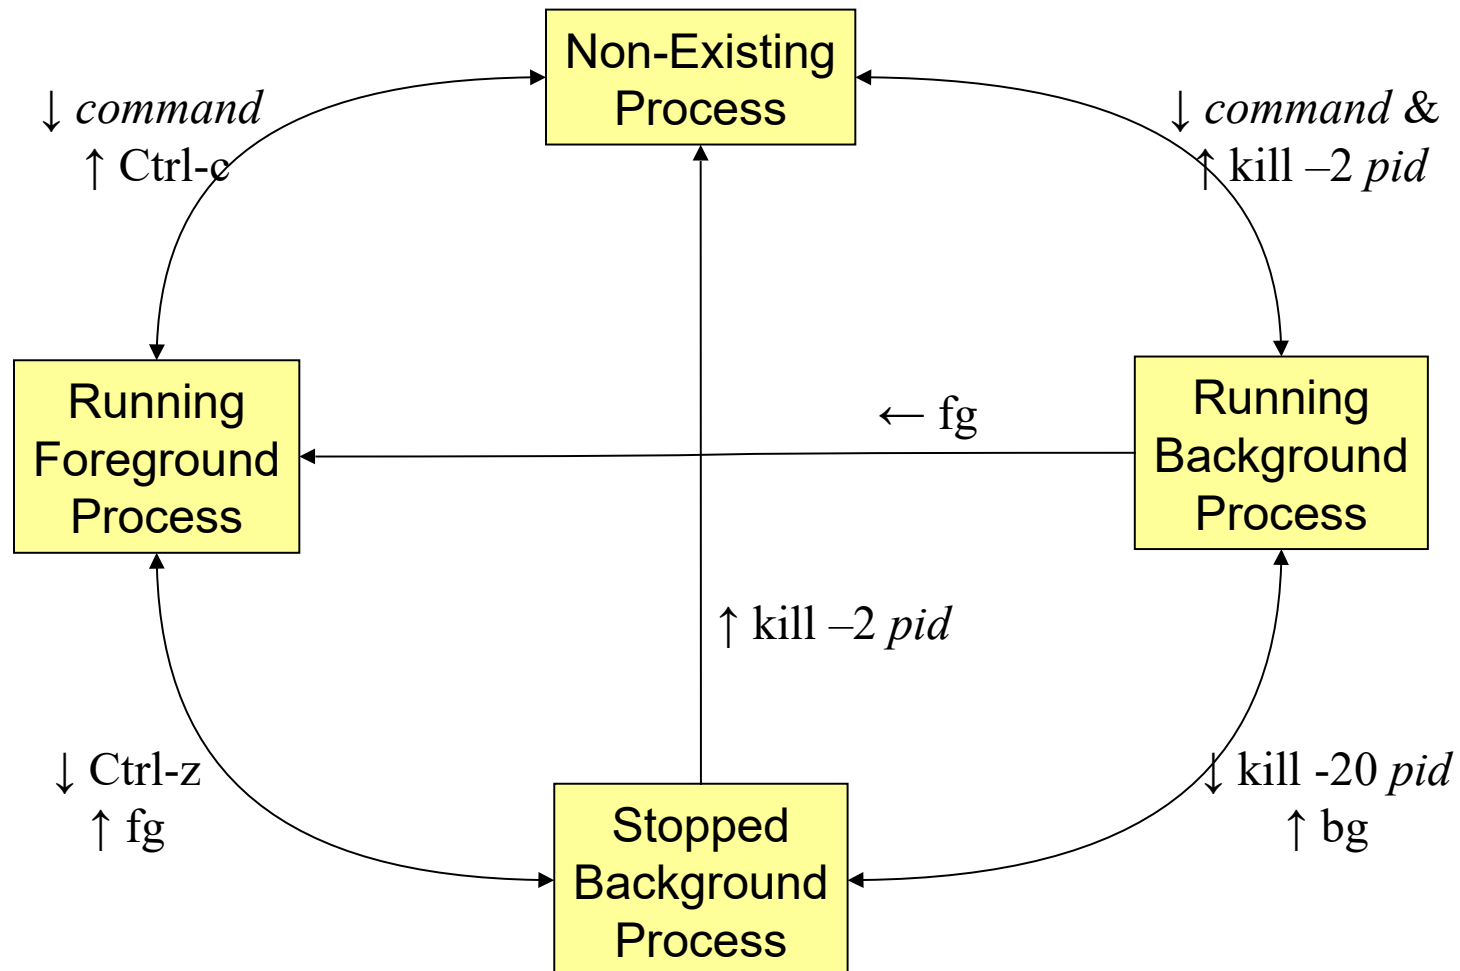
Next lecture: Q2

# Outline



1. **UNIX Process Control**
2. Signals
3. C90 Signal Handling
4. C90 Signal Blocking
5. POSIX Signal Handling/Blocking
6. Conclusion
7. (optional) Alarms and Interval Timers

# UNIX Process Control



# UNIX Process Control



[Demo of UNIX process control using infloop.c]



# Process Control Implementation

Exactly what happens when you:

- Type Ctrl-c?
  - Keyboard sends hardware interrupt
  - Hardware interrupt is handled by OS
  - OS sends a 2/SIGINT **signal**
- Type Ctrl-z?
  - Keyboard sends hardware interrupt
  - Hardware interrupt is handled by OS
  - OS sends a 20/SIGTSTP **signal**
- Issue a “kill *–sig pid*” command?
  - OS sends a *sig* **signal** to the process whose id is *pid*
- Issue a “fg” or “bg” command?
  - OS sends a 18/SIGCONT **signal** (and does some other things too!) 7

# Outline



1. UNIX Process Control
- 2. Signals**
3. C90 Signal Handling
4. C90 Signal Blocking
5. POSIX Signal Handling/Blocking
6. Conclusion
7. (optional) Alarms and Interval Timers



# Signals



Q1: How does the **OS** communicate to an **application process**?


A1: **Signals**

# Definition of Signal

## Signal: A notification of an event

- Event gains attention of the OS
- OS stops the application process immediately, sending it a signal
- **Signal handler** executes to completion
- Application process resumes where it left off

### Process



# Examples of Signals

## User types Ctrl-c

- Event gains attention of OS
- OS stops the application process immediately, sending it a 2/SIGINT signal
- Signal handler for 2/SIGINT signal executes to completion
  - Default signal handler for 2/SIGINT signal exits process



## Process makes illegal memory reference

- Event gains attention of OS
- OS stops application process immediately, sending it a 11/SIGSEGV signal
- Signal handler for 11/SIGSEGV signal executes to completion
  - Default signal handler for 11/SIGSEGV signal prints “segmentation fault” and exits process

# Sending Signals via Keystrokes



Three signals can be sent from keyboard:

- Ctrl-c → 2/SIGINT signal
  - Default handler exits process
- Ctrl-z → 20/SIGTSTP signal
  - Default handler suspends process
- Ctrl-\ → 3/SIGQUIT signal
  - Default handler exits process



# Sending Signals via Commands

## kill Command

`kill -signal pid`

- Send a signal of type **signal** to the process with id **pid**
- Can specify either signal type name (-SIGINT) or number (-2)
- No signal type name or number specified => sends 15/SIGTERM signal
  - Default 15/SIGTERM handler exits process
- Editorial comment: Better command name would be **sendsig**

## Examples

`kill -2 1234`

`kill -SIGINT 1234`

- Same as pressing Ctrl-c if process 1234 is running in foreground

# Sending Signals via Function Call



## raise()

```
int raise(int iSig);
```

- Commands OS to send a signal of type `iSig` to current process
- Returns 0 to indicate success, non-0 to indicate failure

## Example

```
int ret = raise(SIGINT); /* Process commits suicide. */  
assert(ret != 0);        /* Shouldn't get here. */
```

Note: C90 function

# Sending Signals via Function Call



## `kill()`

```
int kill(pid_t iPid, int iSig);
```

- Sends a `iSig` signal to the process whose id is `iPid`
- Equivalent to `raise(iSig)` when `iPid` is the id of current process
- Editorial comment: Better function name would be `sendsig()`

## Example

```
pid_t iPid = getpid(); /* Process gets its id.*/  
kill(iPid, SIGINT); /* Process sends itself a  
                     SIGINT signal (commits  
                     suicide?) */
```

Note: POSIX (not C90) function



# Signal Handling

Each signal type has a default handler

- Most default handlers exit the process

A program can install its own handler for signals of any type

Exceptions: A program *cannot* install its own handler for signals of type:

- 9/SIGKILL
  - Default handler exits the process
  - Catchable termination signal is 15/SIGTERM
- 19/SIGSTOP
  - Default handler suspends the process
  - Can resume the process with signal 18/SIGCONT
  - Catchable suspension signal is 20/SIGTSTP



# Outline



1. UNIX Process Control
2. Signals
3. **C90 Signal Handling**
4. C90 Signal Blocking
5. POSIX Signal Handling/Blocking
6. Conclusion
7. (optional) Alarms and Interval Timers



# Installing a Signal Handler

## `signal()`

```
sighandler_t signal(int iSig,  
                    sighandler_t pfHandler);
```

- Installs function **pfHandler** as the handler for signals of type **iSig**
- **pfHandler** is a function pointer:  

```
typedef void (*sighandler_t)(int);
```
- Returns the old handler on success, **SIG\_ERR** on error
- After call, **pfHandler** is invoked whenever process receives a signal of type **iSig**



# Installing a Handler Example 1

## Program testsignal.c:

```
#define _GNU_SOURCE /* Use modern handling style */
#include <stdio.h>
#include <assert.h>
#include <signal.h>

static void myHandler(int iSig) {
    printf("In myHandler with argument %d\n", iSig);
}
...
```

# Installing a Handler Example 1 (cont.)



## Program testsignal.c (cont.):

```
...
int main(void) {
    void (*pfRet)(int);
    pfRet = signal(SIGINT, myHandler);
    assert(pfRet != SIG_ERR);

    printf("Entering an infinite loop\n");
    for (;;)
        ;
    return 0;
}
```

# Installing a Handler Example 1 (cont.)



[Demo of testsignal.c]



# Installing a Handler Example 2

## Program testsignalall.c:

```
#define _GNU_SOURCE
#include <stdio.h>
#include <assert.h>
#include <signal.h>

static void myHandler(int iSig) {
    printf("In myHandler with argument %d\n", iSig);
}

...
```

# Installing a Handler Example 2 (cont.)



## Program testsignalall.c (cont.):

```
...
int main(void) {
    void (*pfRet) (int);
    pfRet = signal(SIGHUP, myHandler); /* 1 */
    pfRet = signal(SIGINT, myHandler); /* 2 */
    pfRet = signal(SIGQUIT, myHandler); /* 3 */
    pfRet = signal(SIGILL, myHandler); /* 4 */
    pfRet = signal(SIGTRAP, myHandler); /* 5 */
    pfRet = signal(SIGABRT, myHandler); /* 6 */
    pfRet = signal(SIGBUS, myHandler); /* 7 */
    pfRet = signal(SIGFPE, myHandler); /* 8 */
    pfRet = signal(SIGKILL, myHandler); /* 9 */
    ...
}
```

# Installing a Handler Example 2 (cont.)



## Program testsignalall.c (cont.):

```
...
/* Etc., for every signal. */

printf("Entering an infinite loop\n");
for (;;)
    ;
return 0;
}
```



# Installing a Handler Example 2 (cont.)



[Demo of testsignalall.c]



# Installing a Handler Example 3

## Program generates lots of temporary data

- Stores the data in a temporary file
- Must delete the file before exiting

```
...  
int main(void) {  
    FILE *psFile;  
    psFile = fopen("temp.txt", "w");  
    ...  
    fclose(psFile);  
    remove("temp.txt");  
    return 0;  
}
```



# Example 3 Problem

## What if user types Ctrl-c?

- OS sends a 2/SIGINT signal to the process
- Default handler of 2/SIGINT exits the process

## Problem: The temporary file is not deleted

- Process dies before `remove("tmp.txt")` is executed

## Challenge: Ctrl-c could happen at any time

- Which line of code will be interrupted???

## Solution: Install a signal handler

- Define a “clean up” function to delete the file
- Install the function as a signal handler for 2/SIGINT



# Example 3 Solution

```
...
static FILE *psFile; /* Must be global. */
static void cleanup(int iSig) {
    fclose(psFile);
    remove("tmp.txt");
    exit(EXIT_FAILURE);
}
int main(void) {
    void (*pfRet)(int);
    psFile = fopen("temp.txt", "w");
    pfRet = signal(SIGINT, cleanup);
    ...
    raise(SIGINT);
    return 0; /* Never get here. */
}
```



# Predefined Signal Handler: SIG\_IGN

Pre-defined signal handler: SIG\_IGN

Can install to ignore signals

```
int main(void) {  
    void (*pfRet)(int);  
    pfRet = signal(SIGINT, SIG_IGN);  
    ...  
}
```

Subsequently, process will ignore SIGINT signals



# Predefined Signal Handler: SIG\_DFL

Pre-defined signal handler: SIG\_DFL

Can install to restore default signal handler

```
int main(void) {  
    void (*pfRet)(int);  
    pfRet = signal(SIGINT, somehandler);  
    ...  
    pfRet = signal(SIGINT, SIG_DFL);  
    ...  
}
```

Subsequently, process will handle 2/SIGINT signals using the default handler for 2/SIGINT signals

# Outline



1. UNIX Process Control
2. Signals
3. C90 Signal Handling
- 4. C90 Signal Blocking**
5. POSIX Signal Handling/Blocking
6. Conclusion
7. (optional) Alarms and Interval Timers

# Race Conditions in Signal Handlers



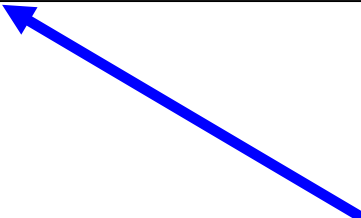
A **race condition** is a flaw in a program whereby the correctness of the program is critically dependent on the sequence or timing of other events.

Race conditions can occur in signal handlers...



# Race Condition Example

```
void addSalaryToSavings(int iSig) {  
    int iTemp;  
    iTemp = iSavingsBalance;  
    iTemp += iMonthlySalary;  
    iSavingsBalance = iTemp;  
}
```

A blue arrow points from the text below to the function definition above.

Handler for hypothetical  
“update monthly salary” signal



# Race Condition Example (cont.)

(1) Signal arrives; handler begins executing

```
void addSalaryToSavings(int iSig) {  
    int iTemp;  
    ↓ iTemp = iSavingsBalance;    2000  
    iTemp += iMonthlySalary;  
    iSavingsBalance = iTemp;  
}
```

# Race Condition Example (cont.)

(2) Another signal arrives; first instance of handler is interrupted; second instance of handler begins executing

```
void addSalaryToSavings(int iSig) {  
    int iTemp;  
    ↓ iTemp = iSavingsBalance;    2000  
    iTemp += iMonthlySalary;  
    iSavingsBalance = iTemp;  
}
```

```
void addSalaryToSavings(int iSig) {  
    int iTemp;  
    ↓ iTemp = iSavingsBalance;    2000  
    iTemp += iMonthlySalary;  
    iSavingsBalance = iTemp;  
}
```

# Race Condition Example (cont.)

## (3) Second instance executes to completion

```
void addSalaryToSavings(int iSig) {  
    ↓ int iTemp;  
    ↓ iTemp = iSavingsBalance;    2000  
    iTemp += iMonthlySalary;  
    iSavingsBalance = iTemp;  
}
```

```
void addSalaryToSavings(int iSig) {  
    int iTemp;  
    iTemp = iSavingsBalance;    2000  
    iTemp += iMonthlySalary;    2050  
    ↓ iSavingsBalance = iTemp;    2050  
}
```



# Race Condition Example (cont.)

(4) Control returns to first instance, which executes to completion

```
void addSalaryToSavings(int iSig) {  
    int iTemp;  
    ↓ iTemp = iSavingsBalance;      2000  
    ↓ iTemp += iMonthlySalary;      2050  
    ↓ iSavingsBalance = iTemp;      2050  
}
```

Lost 50 !!!



# Blocking Signals in Handlers

## Blocking signals

- To **block** a signal is to **queue** it for delivery at a later time

## Why block signals when handler is executing?

- Avoid race conditions when another signal of type x occurs while the handler for type x is executing

## How to block signals when handler is executing?

- **Automatic** during execution of signal handler!!!
- Previous sequence **cannot happen!!!**
- While executing a handler for a signal of type x, all signals of type x are blocked
- When/if signal handler returns, block is removed

# Race Conditions in General

Race conditions can occur elsewhere too

```
int iFlag = 0;

void myHandler(int iSig) {
    iFlag = 1;
}

int main(void) {
    if (iFlag == 0) {
        /* Do something */
    }
}
```

Problem: myflag might become 1 just after the comparison!

Must make sure that **critical sections** of code are not interrupted



# Blocking Signals in General

## How to block signals in general?

- Not possible in C90
- Possible using POSIX functions...



# Outline



1. UNIX Process Control
2. Signals
3. C90 Signal Handling
4. C90 Signal Blocking
- 5. POSIX Signal Handling/Blocking**
6. Conclusion
7. (optional) Alarms and Interval Timers



# POSIX Signal Handling

## C90 standard

- Defines `signal()` and `raise()` functions
  - Work across all systems (UNIX, LINUX, Windows), but...
  - Work **differently** across some systems!!!
    - On some systems, signals are blocked during execution of handler for that type of signal -- but not so on other (older) systems
    - On some (older) systems, handler installation for signals of type x is cancelled after first signal of type x is received; must reinstall the handler -- but not so on other systems
- Does not provide mechanism to block signals in general



# POSIX Signal Handling

## POSIX standard

- Defines `kill()`, `sigprocmask()`, and `sigaction()` functions
  - Work the same across all POSIX-compliant UNIX systems (Linux, Solaris, etc.), but...
  - Do not work on non-UNIX systems (e.g. Windows)
- Provides mechanism to block signals in general



# Blocking Signals in General

Each process has a signal mask in the kernel

- OS uses the mask to decide which signals to deliver
- User program can modify mask with `sigprocmask()`

**`sigprocmask()`**

```
int sigprocmask(int iHow,  
                const sigset_t *psSet,  
                sigset_t *psOldSet);
```

- **`psSet`**: Pointer to a signal set
- **`psOldSet`**: (Irrelevant for our purposes)
- **`iHow`**: How to modify the signal mask
  - **`SIG_BLOCK`**: Add **`psSet`** to the current mask
  - **`SIG_UNBLOCK`**: Remove **`psSet`** from the current mask
  - **`SIG_SETMASK`**: Install **`psSet`** as the signal mask
- Returns 0 iff successful

Functions for constructing signal sets

- **`sigemptyset()`**, **`sigaddset()`**, ...

Note: No parallel function in C90



# Blocking Signals Example

```
sigset_t sSet;
int main(void) {
    int iRet;
    sigemptyset(&sSet);
    sigaddset(&sSet, SIGINT);
    iRet = sigprocmask(SIG_BLOCK, &sSet, NULL);
    assert(iRet == 0);
    if (iFlag == 0) {
        /* Do something */
    }
    iRet = sigprocmask(SIG_UNBLOCK, &sSet, NULL);
    assert(iRet == 0);
    ...
}
```



# Blocking Signals in Handlers

Signals of type  $x$  automatically are blocked when executing handler for signals of type  $x$

Additional signal types to be blocked can be defined at time of handler installation...



# Installing a Signal Handler

## `sigaction()`

```
int sigaction(int iSig,  
              const struct sigaction *psAction,  
              struct sigaction *psOldAction);
```

- `iSig`: The type of signal to be affected
- `psAction`: Pointer to a structure containing instructions on how to handle signals of type `iSig`, including signal handler name and which signal types should be blocked
- `psOldAction`: (Irrelevant for our purposes)
- Installs an appropriate handler
- Automatically blocks signals of type `iSig`
- Returns 0 iff successful

Note: More powerful than C90 `signal()`



# Installing a Handler Example

Program testsigaction.c:

```
#define _GNU_SOURCE
#include <stdio.h>
#include <stdlib.h>
#include <signal.h>

static void myHandler(int iSig) {
    printf("In myHandler with argument %d\n", iSig);
}

...
```





# Installing a Handler Example (cont.)

Program testsigaction.c (cont.):

```
...
int main(void) {
    int iRet;
    struct sigaction sAction;
    sAction.sa_flags = 0;
    sAction.sa_handler = myHandler;
    sigemptyset(&sAction.sa_mask);
    iRet = sigaction(SIGINT, &sAction, NULL);
    assert(iRet == 0);

    printf("Entering an infinite loop\n");
    for (;;)
        ;
    return 0;
}
```

# Installing a Handler Example (cont.)



[Demo of testsigaction.c]

# Outline



1. UNIX Process Control
2. Signals
3. C90 Signal Handling
4. C90 Signal Blocking
5. POSIX Signal Handling/Blocking
- 6. Conclusion**
7. (optional) Alarms and Interval Timers



# Predefined Signals

List of the predefined signals:

```
$ kill -l
```

1) SIGHUP	2) <b>SIGINT</b>	3) <b>SIGQUIT</b>	4) SIGILL
5) SIGTRAP	6) SIGABRT	7) SIGBUS	8) SIGFPE
9) <b>SIGKILL</b>	10) SIGUSR1	11) <b>SIGSEGV</b>	12) SIGUSR2
13) SIGPIPE	14) <b>SIGALRM</b>	15) <b>SIGTERM</b>	17) SIGCHLD
18) <b>SIGCONT</b>	19) <b>SIGSTOP</b>	20) <b>SIGTSTP</b>	21) SIGTTIN
22) SIGTTOU	23) SIGURG	24) SIGXCPU	25) SIGXFSZ
26) SIGVTALRM	27) <b>SIGPROF</b>	28) SIGWINCH	29) SIGIO
30) SIGPWR	31) SIGSYS	34) SIGRTMIN	35) SIGRTMIN+1
36) SIGRTMIN+2	37) SIGRTMIN+3	38) SIGRTMIN+4	39) SIGRTMIN+5
40) SIGRTMIN+6	41) SIGRTMIN+7	42) SIGRTMIN+8	43) SIGRTMIN+9
44) SIGRTMIN+10	45) SIGRTMIN+11	46) SIGRTMIN+12	47) SIGRTMIN+13
48) SIGRTMIN+14	49) SIGRTMIN+15	50) SIGRTMAX-14	51) SIGRTMAX-13
52) SIGRTMAX-12	53) SIGRTMAX-11	54) SIGRTMAX-10	55) SIGRTMAX-9
56) SIGRTMAX-8	57) SIGRTMAX-7	58) SIGRTMAX-6	59) SIGRTMAX-5
60) SIGRTMAX-4	61) SIGRTMAX-3	62) SIGRTMAX-2	63) SIGRTMAX-1
64) SIGRTMAX			

Applications can define their own signals

- An application can define signals with unused values

# Summary



## Signals

- A **signal** is an asynchronous event mechanism
- C90 **raise()** or POSIX **kill()** sends a signal
- C90 **signal()** or POSIX **sigaction()** installs a signal handler
  - Most predefined signals are “catchable”
- Beware of race conditions
- Signals of type **x** automatically are blocked while handler for type **x** signals is running
- POSIX **sigprocmask()** blocks signals in any critical section of code

# Summary



Q: How does the OS communicate to application programs?

A: Signals

For more information:

Bryant & O'Hallaron, *Computer Systems: A Programmer's Perspective*, Chapter 8

# Outline



1. UNIX Process Control
2. Signals
3. C90 Signal Handling
4. C90 Signal Blocking
5. POSIX Signal Handling/Blocking
6. Conclusion
7. **(optional) Alarms and Interval Timers**

# Alarms

## `alarm()`

`unsigned int alarm(unsigned int uiSec);`

- Sends 14/SIGALRM signal after `uiSec` seconds
- Cancels pending alarm if `uiSec` is 0
- Uses **real time**, alias **wall-clock time**
  - Time spent executing other processes counts
  - Time spent waiting for user input counts
- Return value is meaningless

Used to implement time-outs







# Alarm Example 1

Program testalarm.c:

```
#define _GNU_SOURCE
#include <stdio.h>
#include <assert.h>
#include <signal.h>
#include <unistd.h>

static void myHandler(int iSig) {
    printf("In myHandler with argument %d\n", iSig);

    /* Set another alarm. */
    alarm(2);
}

...
```



# Alarm Example 1 (cont.)

## Program testalarm.c (cont.):

```
...
int main(void)
{
    void (*pfRet) (int) ;
    sigset_t sSet;
    int iRet;

    /* Make sure that SIGALRM is not blocked. */
    sigemptyset(&sSet);
    sigaddset(&sSet, SIGALRM);
    iRet = sigprocmask(SIG_UNBLOCK, &sSet, NULL);
    assert(iRet == 0);

    pfRet = signal(SIGALRM, myHandler);
    assert(pfRet != SIG_ERR);
    ...
}
```

# Alarm Example 1 (cont.)

Program testalarm.c (cont.):

```
...  
  
/* Set an alarm. */  
alarm(2) ;  
  
printf("Entering an infinite loop\n");  
for (;;)   
    ;  
  
return 0;  
}
```

# Alarm Example 1 (cont.)



[Demo of testalarm.c]

# Alarm Example 2



Program `testalarmtimeout.c`:

```
#define _GNU_SOURCE
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include <signal.h>
#include <unistd.h>

static void myHandler(int iSig)
{
    printf("\nSorry.  You took too long.\n");
    exit(EXIT_FAILURE);
}
```



# Alarm Example 2 (cont.)

## Program testalarmtimeout.c (cont.):

```
int main(void) {
    int i;
    void (*pfRet)(int);
    sigset_t sSet;
    int iRet;

    /* Make sure that SIGALRM is not blocked. */
    sigemptyset(&sSet);
    sigaddset(&sSet, SIGALRM);
    iRet = sigprocmask(SIG_UNBLOCK, &sSet, NULL);
    assert(iRet == 0);
    ...
}
```



# Alarm Example 2 (cont.)

Program testalarmtimeout.c (cont.):

```
...  
pfRet = signal(SIGALRM, myHandler);  
assert(pfRet != SIG_ERR);  
  
printf("Enter a number:  ");  
alarm(5);  
scanf("%d", &i);  
alarm(0);  
  
printf("You entered the number %d.\n", i);  
return 0;  
}
```

# Alarm Example 2 (cont.)



[Demo of testalarmtimeout.c]



# Interval Timers

## `setitimer()`

```
int setitimer(int iWhich,  
              const struct itimerval *psValue,  
              struct itimerval *psOldValue);
```

- Sends 27/SIGPROF signal continually
- Timing is specified by **psValue**
- **psOldValue** is irrelevant for our purposes
- Uses **virtual time**, alias **CPU time**
  - Time spent executing other processes does not count
  - Time spent waiting for user input does not count
- Returns 0 iff successful

Used by execution profilers

# Interval Timer Example

Program testtimer.c:

```
#define _GNU_SOURCE
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include <signal.h>
#include <sys/time.h>

static void myHandler(int iSig) {
    printf("In myHandler with argument %d\n", iSig);
}

...
```

# Interval Timer Example (cont.)

## Program testtimer.c (cont.):

```
...  
int main(void)  
{  
    int iRet;  
    void (*pfRet) (int);  
    struct itimerval sTimer;  
  
    pfRet = signal(SIGPROF, myHandler);  
    assert(pfRet != SIG_ERR);  
    ...  
}
```



# Interval Timer Example (cont.)

## Program testitimer.c (cont.):

```
...
/* Send first signal in 1 second, 0 microseconds. */
sTimer.it_value.tv_sec = 1;
sTimer.it_value.tv_usec = 0;

/* Send subsequent signals in 1 second,
   0 microseconds intervals. */
sTimer.it_interval.tv_sec = 1;
sTimer.it_interval.tv_usec = 0;

iRet = setitimer(ITIMER_PROF, &sTimer, NULL);
assert(iRet != -1);

printf("Entering an infinite loop\n");
for (;;)
    ;
return 0;
}
```

# Interval Timer Example (cont.)



[Demo of testitimer.c]

# Summary



## Alarms

- Call `alarm()` to deliver 14/SIGALRM signals in real/wall-clock time
- Alarms can be used to implement time-outs

## Interval Timers

- Call `setitimer()` to deliver 27/SIGPROF signals in virtual/CPU time
- Interval timers are used by execution profilers